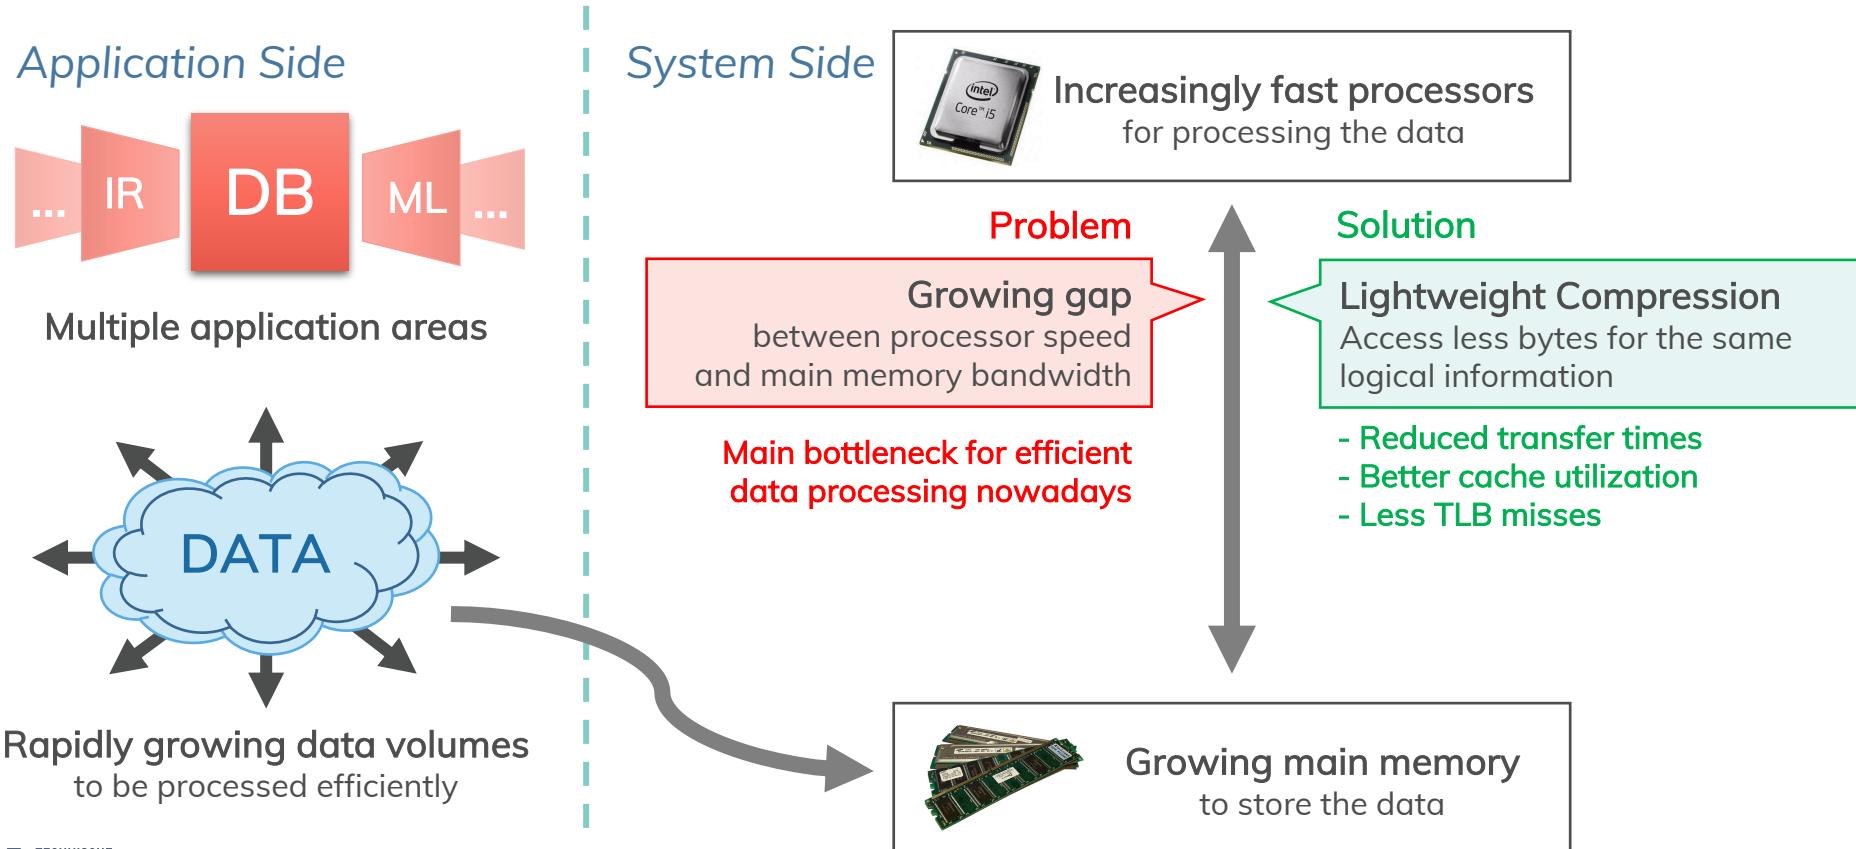


Conflict Detection-based Run-Length Encoding – AVX-512 CD Instruction Set in Action

Annett Ungethüm, Johannes Pietrzyk, Patrick Damme, Dirk Habich, Wolfgang Lehner

HardBD & Active'18 Workshop in Paris, France on April 16, 2018

Challenges for Data Processing Nowadays



Lightweight Compression Techniques

Technique = abstract idea of how compression works

RLE	DELTA	FOR	DICT	NS
Run Length Encoding Replace run by value & length	Differential Coding Replace data elem. by difference to predecessor	Frame-of-Reference Replace data elem. by difference to reference value	Dictionary Coding Replace data elem. by 0-based key in dictionary	Null Suppression Eliminate leading zeroes in binary representation
1200 1200	1000 1000	1200 200	1000 0	00...00 1011
1200 4	1100 100	1100 100	1200 1	↓
1200 300	1150 50	1000 0	1000 0	1011
1200 2	1350 200	1050 50	1050 2	
300	1355 5			
300				

Vectorization is crucial from performance perspective

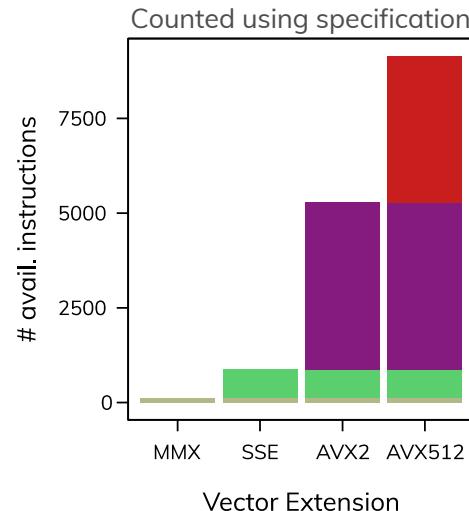
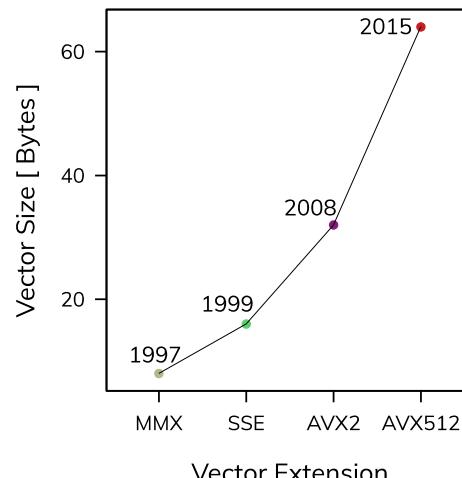
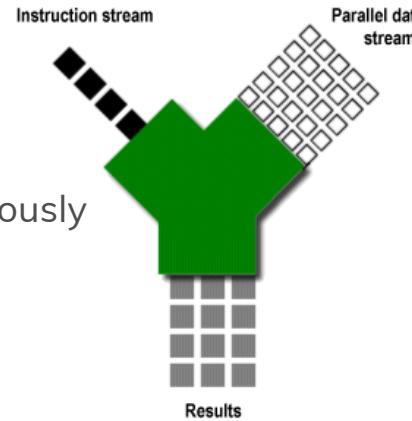
Vectorization using SIMD

Single Instruction Multiple Data (SIMD)

- same instruction on multiple data points simultaneously

Development of Intel's SIMD Extension

- Trend to larger vector registers
 - 128-bit (SSE)
 - 256-bit (AVX and AVX2)
 - 512-bit (AVX-512)
- Trend to more instructions



Vectorization and Lightweight Data Compression

Most algorithms have been proposed for 128-bit SIMD registers

- Processing 4 elements (32 bit integers) at one

Example Run-Length Encoding

- View subsequent occurrences of the same value as a run
- Each run representable by its value and length → just two integers

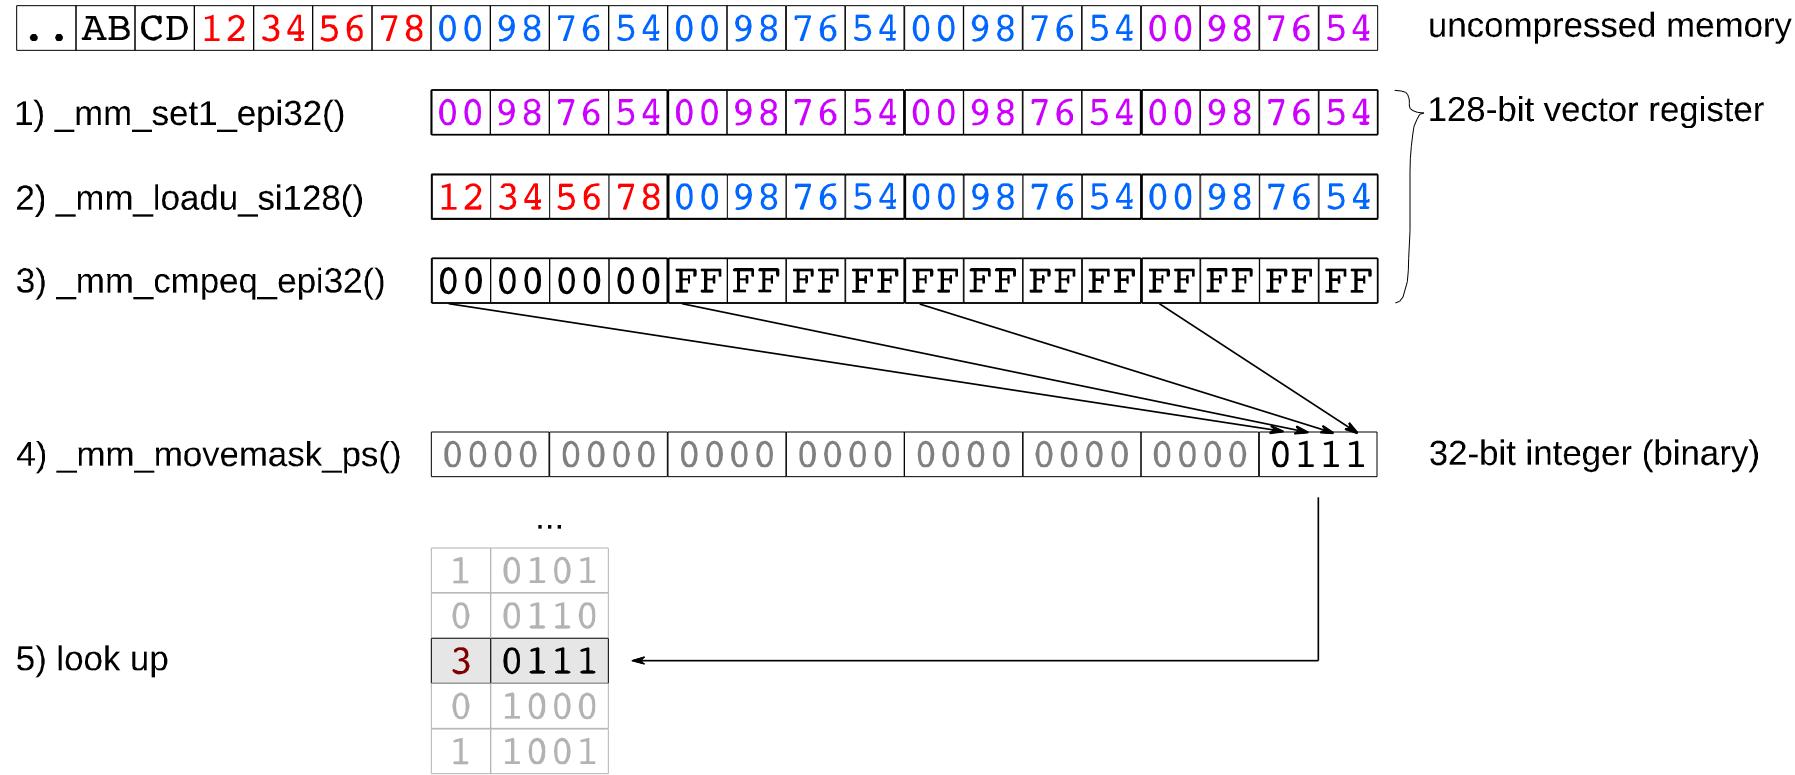
RLE-SIMD

- Uses SIMD instructions to parallelize comparisons

uncompressed				Rle
00	98	76	54	00 98 76 54
00	98	76	54	00 00 00 04
00	98	76	54	12 34 56 78
00	98	76	54	00 00 00 01
12	34	56	78	00 00 AB CD
00	00	AB	CD	00 00 AB CD
00	00	AB	CD	00 00 00 03



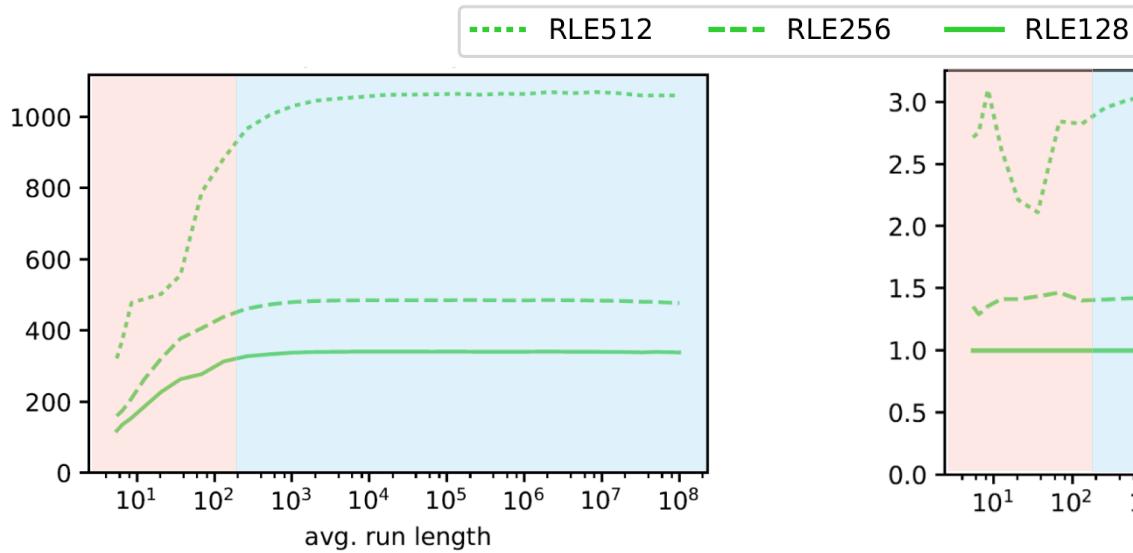
RLE-SIMD: Compression



Evaluation using Different Vector Sizes

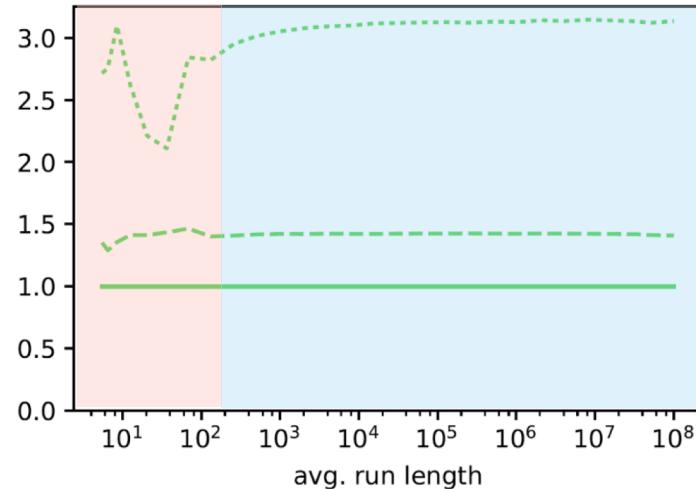
Compression Speed

- Measured in million integers per second (mis)



Speedup

- Compared to baseline of 128-bit



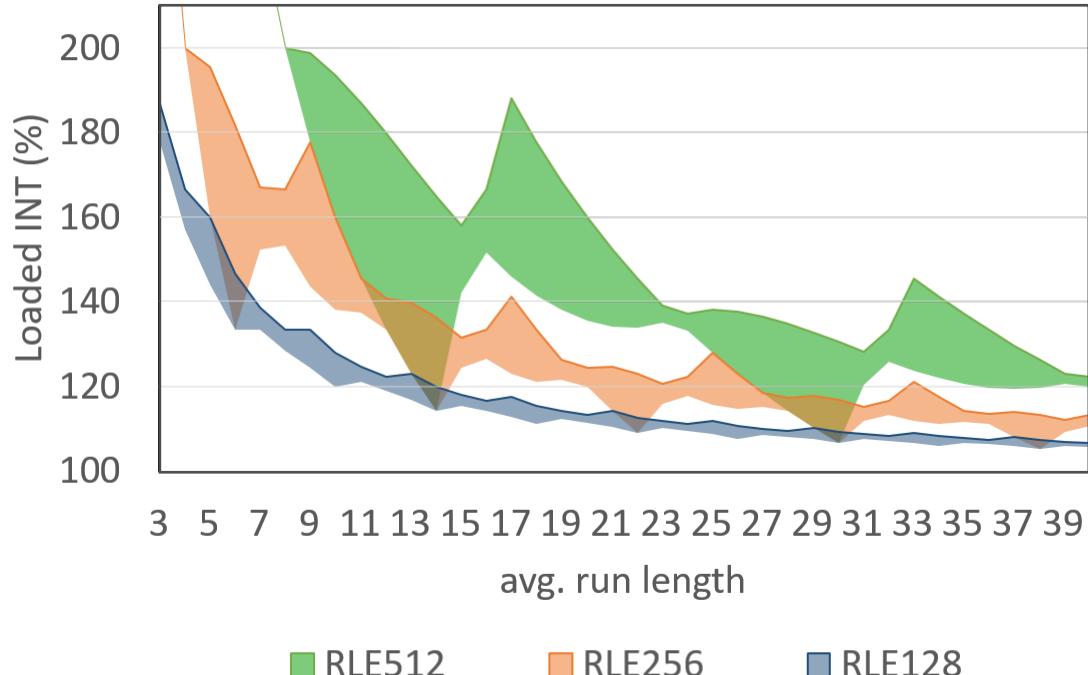
non-well performing area

well-performing area

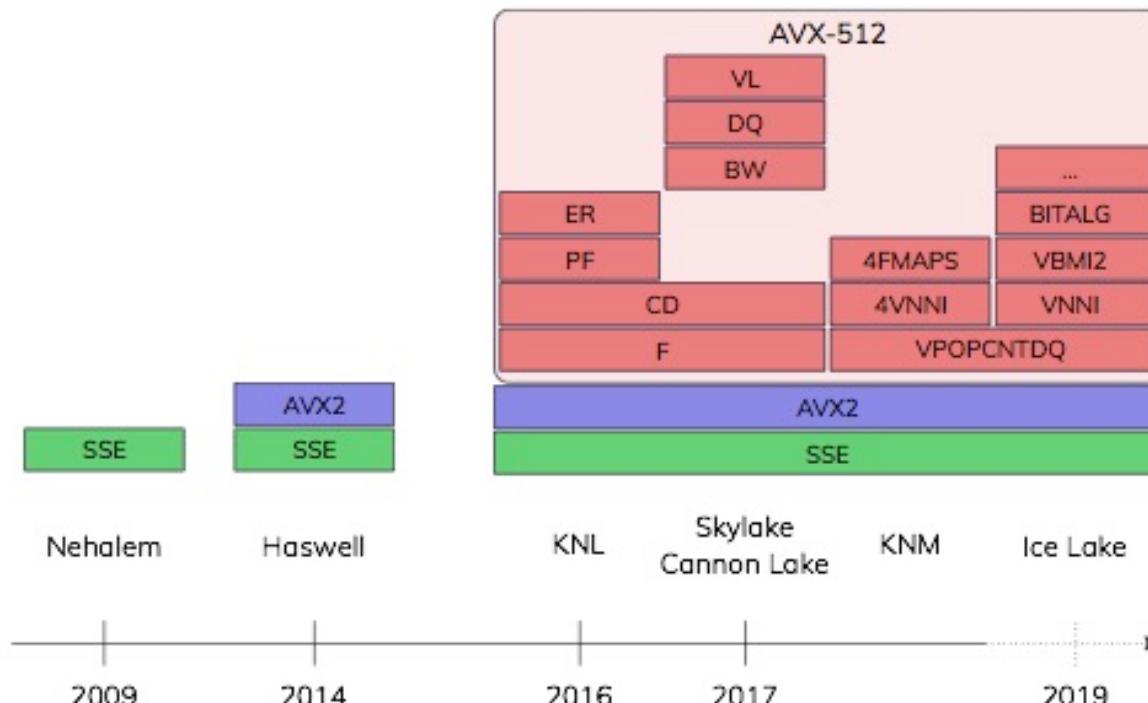
Non-Well Performing Area

Reasons

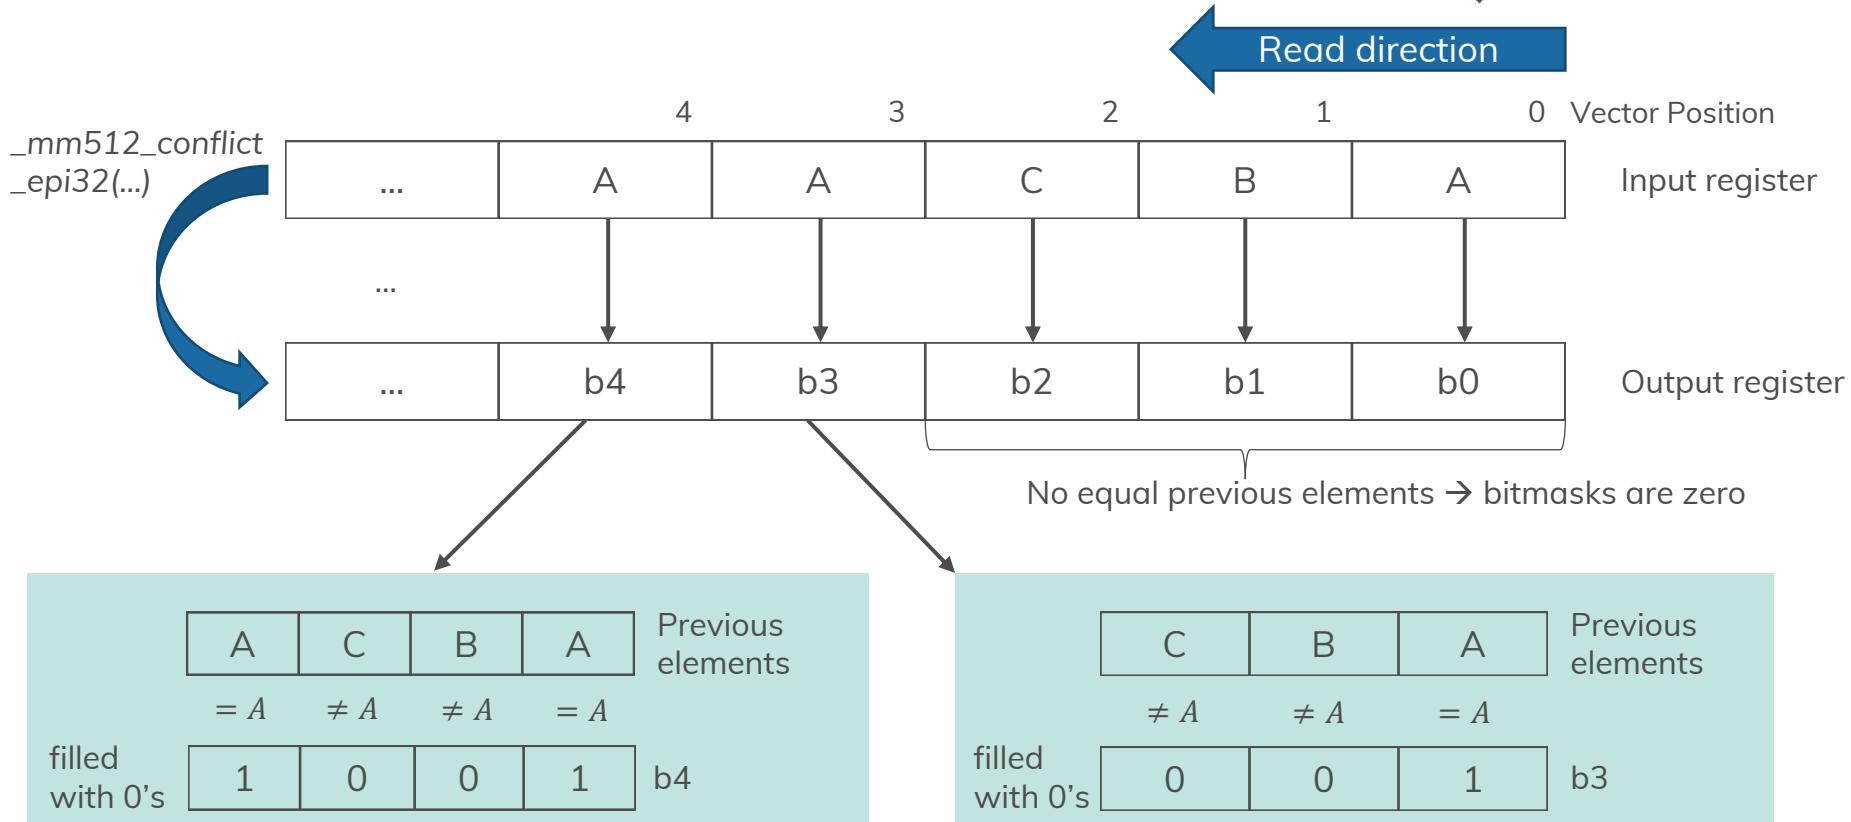
- For large run lengths, the number of loaded integers approaches more or less 100%, i.e. every value is only processed once.
- RLE vectorization uses a significantly higher number of load operations for sequences with short runs.
- The redundant processing dramatically increases with increasing vector widths.



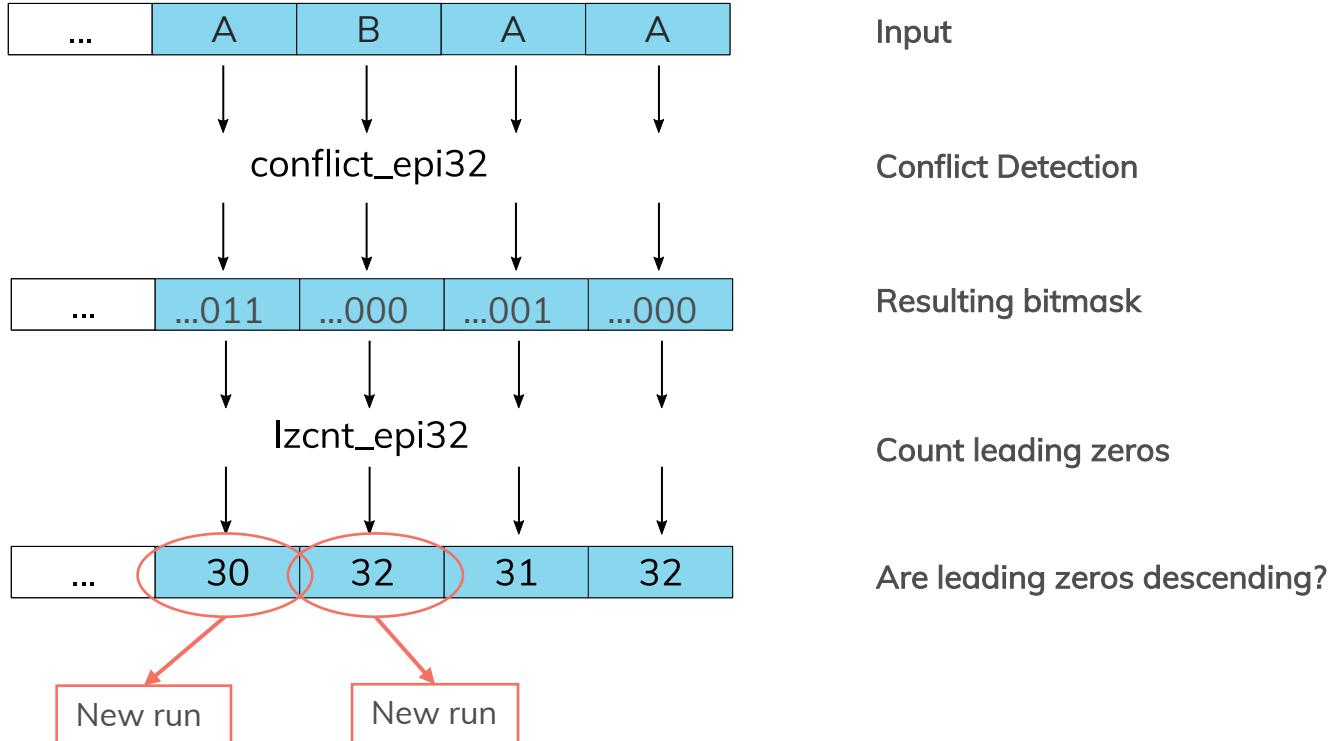
SIMD – New Instruction Sets



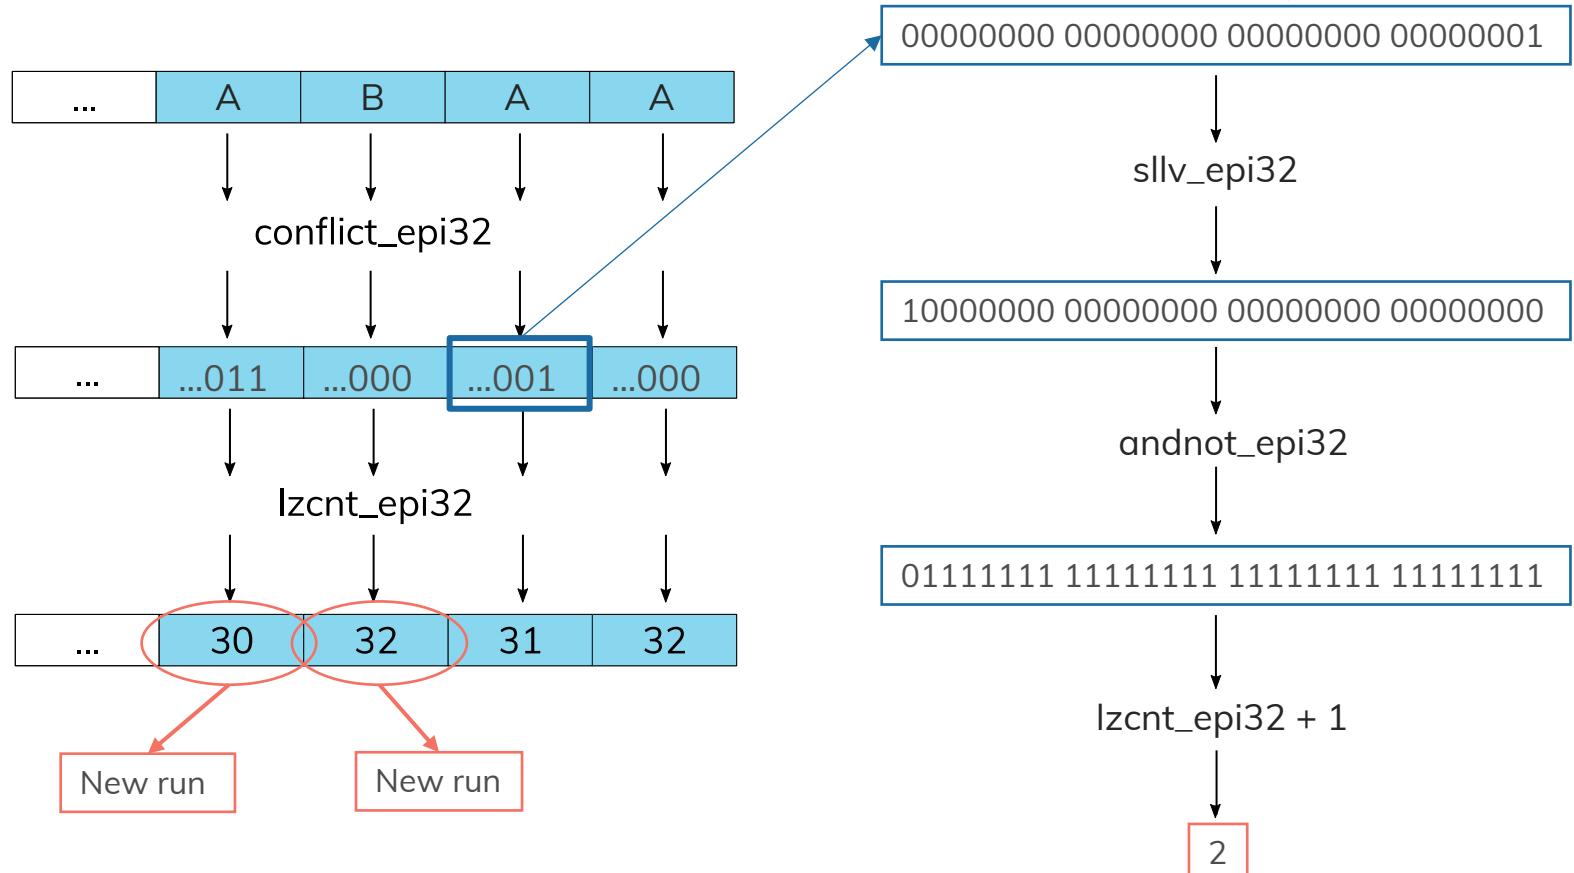
Conflict Detection using AVX512CD



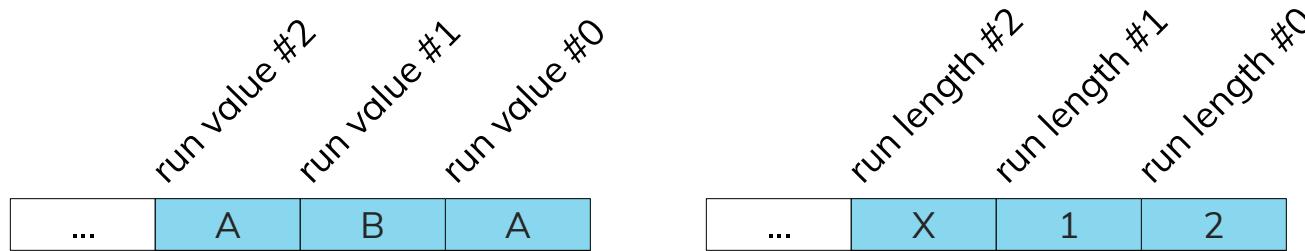
Step 1: Run Detection



Step 2: Run Length Detection



Step 3: Storing



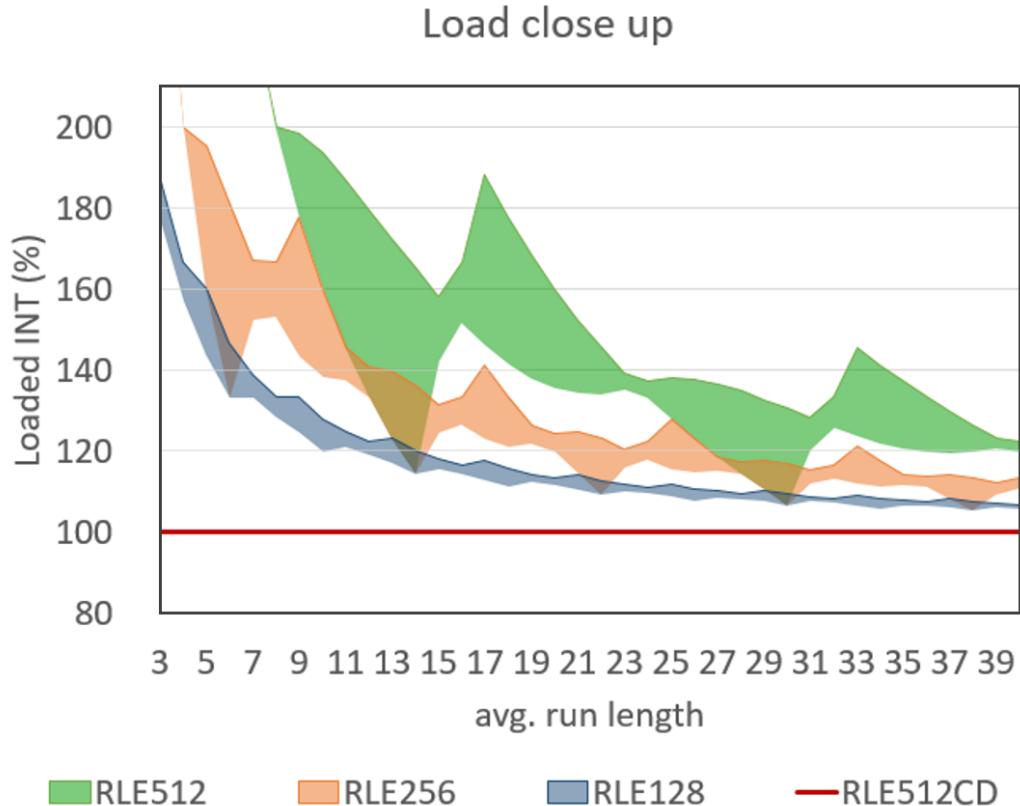
Scatter (RLE512CD)

- Classical storage layout: (run value, run length)-pair
- Independent of vector width

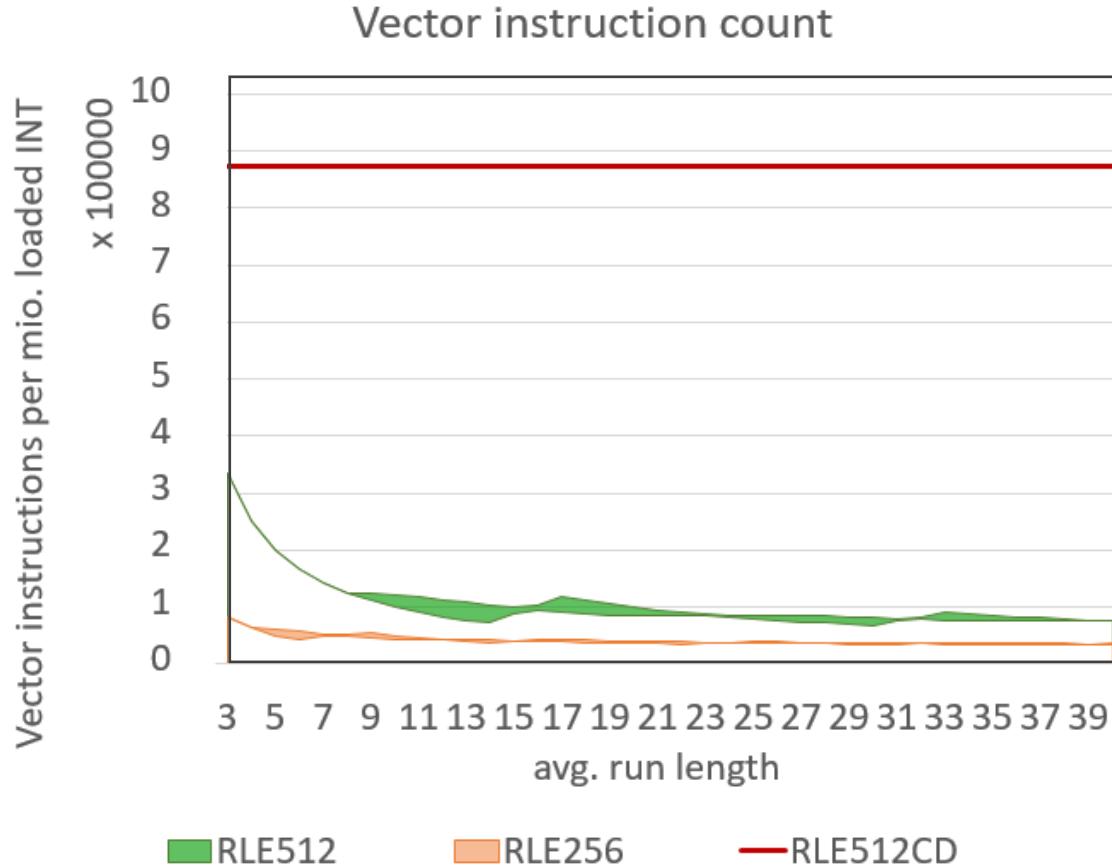
Continuous (RLE512CDAligned)

- Vector wise
- Run length and run value to different memory locations

Evaluation – Load Instructions



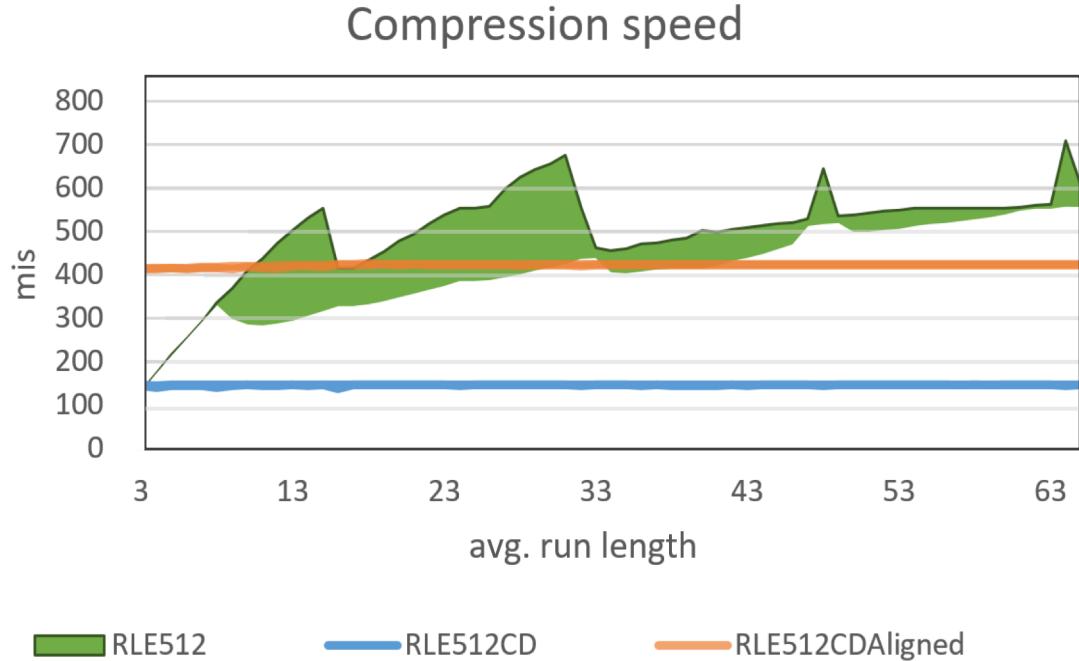
Evaluation- Vector Instructions



Evaluation

Runtime Comparison

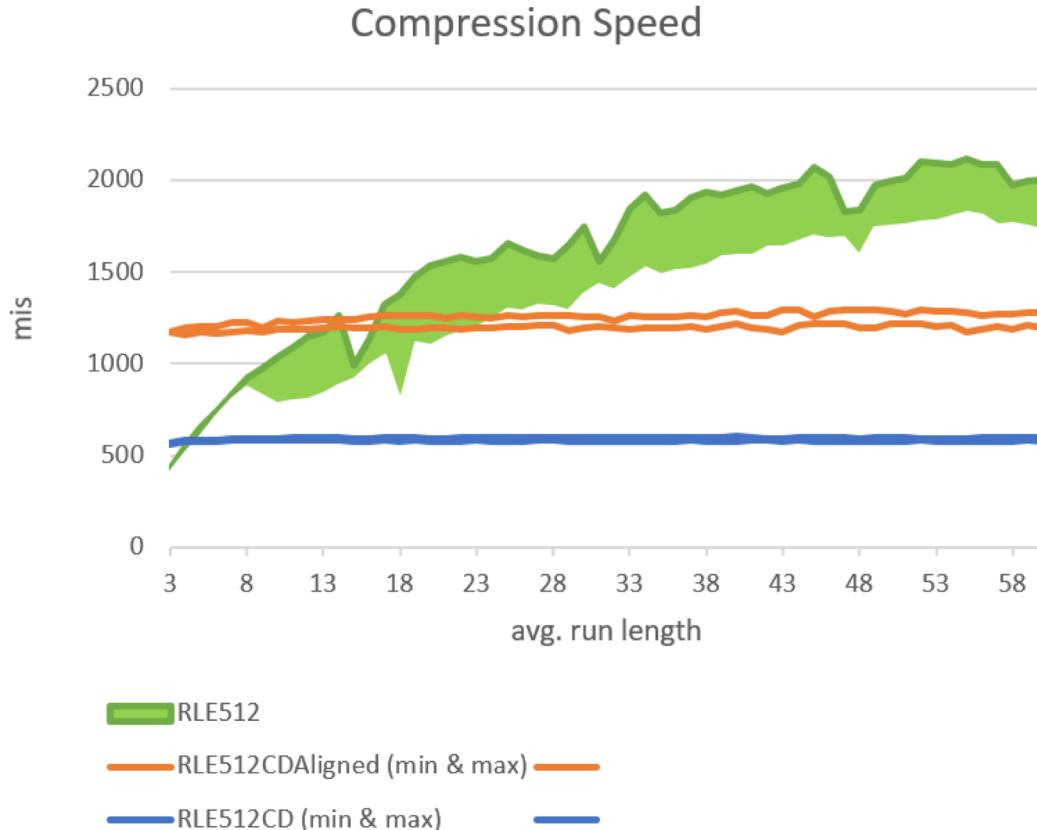
- Intel Xeon Phi Knights Landing Processor
- RLE512CD (Aligned) outperforms state-of-the-art for small average run lengths



Evaluation

Runtime Comparison

- Intel Xeon 6130 Processor
- Similar results



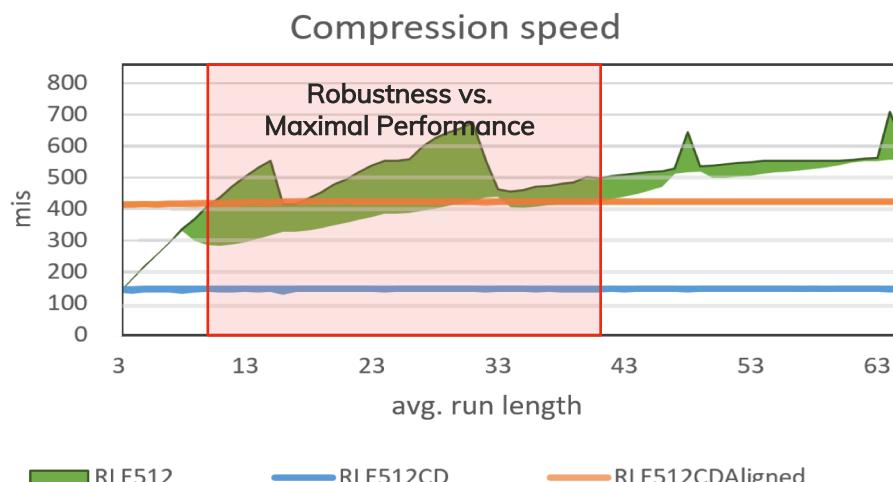
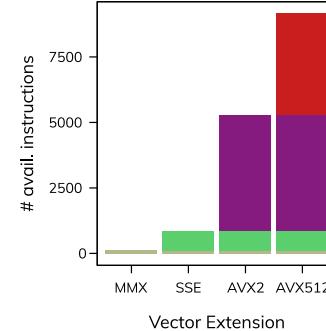
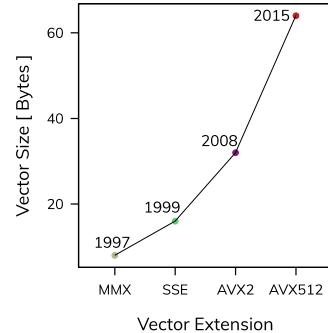
Summary

Development of Intel's SIMD Extension

- Trend to larger vector registers
 - 128-bit (SSE)
 - 256-bit (AVX and AVX2)
 - 512-bit (AVX-512)
- Trend to more instructions

Run Length Encoding

- Proposed novel implementation using AVX512-CD functionality



Conflict Detection-based Run-Length Encoding – AVX-512 CD Instruction Set in Action

Annett Ungethüm, Johannes Pietrzyk, Patrick Damme, Dirk Habich, Wolfgang Lehner

HardBD & Active'18 Workshop in Paris, France on April 16, 2018